

Name: Aarya Admane  
Rollno: 22630  
Div: B Batch: B2

```
import numpy as np

# Sigmoid activation function
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

# Derivative of sigmoid for backpropagation
def sigmoid_derivative(x):
    return x * (1 - x)

# Mean Squared Error Loss function
def mse_loss(y_true, y_pred):
    return np.mean((y_true - y_pred) ** 2)

# Initialize neural network parameters
np.random.seed(42)
input_size = 2    # Number of input neurons
hidden_size = 3   # Number of neurons in hidden layer
output_size = 1   # Number of output neurons

# Initialize weights and biases randomly
W1 = np.random.randn(input_size, hidden_size) # Input to hidden layer
weights
b1 = np.random.randn(1, hidden_size)          # Bias for hidden layer
W2 = np.random.randn(hidden_size, output_size) # Hidden to output
weights
b2 = np.random.randn(1, output_size)           # Bias for output layer

# Training data (XOR problem)
X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]]) # Inputs
y = np.array([[0], [1], [1], [0]])             # Expected output

# Training parameters
epochs = 10000 # Number of training iterations
learning_rate = 0.1 # Learning rate

# Training loop
for epoch in range(epochs):
    # Forward Propagation
    z1 = np.dot(X, W1) + b1 # Linear transformation (Input → Hidden)
    a1 = sigmoid(z1)        # Activation function (Hidden layer)
    z2 = np.dot(a1, W2) + b2 # Linear transformation (Hidden →
```

```

Output)
    a2 = sigmoid(z2)                # Activation function (Output layer)

    # Compute loss
    loss = mse_loss(y, a2)

    # Backpropagation (Calculating Gradients)

    d_a2 = (a2 - y) # Loss gradient
    d_z2 = d_a2 * sigmoid_derivative(a2) # Backprop through activation
    d_W2 = np.dot(a1.T, d_z2) # Weight gradient (hidden → output)
    d_b2 = np.sum(d_z2, axis=0, keepdims=True) # Bias gradient (output)

    d_a1 = np.dot(d_z2, W2.T) # Backprop to hidden layer
    d_z1 = d_a1 * sigmoid_derivative(a1) # Backprop through activation
    d_W1 = np.dot(X.T, d_z1) # Weight gradient (input → hidden)
    d_b1 = np.sum(d_z1, axis=0, keepdims=True) # Bias gradient (hidden)

    # Update weights and biases
    W2 -= learning_rate * d_W2
    b2 -= learning_rate * d_b2
    W1 -= learning_rate * d_W1
    b1 -= learning_rate * d_b1

    # Print loss every 1000 epochs
    if epoch % 1000 == 0:
        print(f"Epoch {epoch}, Loss: {loss:.4f}")

# Final Predictions
print("\nFinal Predictions:")
print(sigmoid(np.dot(sigmoid(np.dot(X, W1) + b1), W2) + b2))

```

Final Predictions:

```

[[0.54553169]
 [0.57471011]
 [0.53845304]
 [0.56280605]]

```